Sumiksh Trehan . Jarvis Consulting

Sumiksh is dynamic and results-oriented Power BI Developer and Data Analyst with a strong technical foundation in Software Development. Expert in transforming raw architectural and testing data into actionable insights by designing high-impact dashboards and engineering complex DAX queries to track critical metrics like POD velocity and defect density.

- Expertise in DAX & Data Architecture: He engineers sophisticated Power BI solutions using complex DAX measures and Star Schema modeling to deliver high-performance analytics, such as his ticketing system dashboard and U.S. demographic projects.
- Domain-Specific QA Intelligence: Leveraging deep STLC expertise from major government programs, he architects reporting suites that track critical metrics like defect density and test execution health to provide end-to-end project visibility.
- Agile Leadership & PM Mindset: Drawing on his experience in banking and government, he bridges the gap between technical data and delivery by facilitating stand-ups, removing POD blockers, and optimizing team velocity.
- Proven Record of Optimization: He has a demonstrated history of driving operational efficiency, notably increasing reporting speeds by 20% through strategic automation and streamlined data workflows.

Skills

Proficient: Power BI, Tableau, JavaScript (Next JS, React, Node.JS), Linux/Bash, RDBMS/SQL, Agile/Scrum, Git, GCP, Docker

Competent: Next.js, React, Azure(Cloud, Data Factory), Kotlin/Android Development, ETL Pipelines(Python, SSIS)

Familiar: AWS, Azure, C#, C, C++, Perl, Cobol

Jarvis Projects

Project source code: https://github.com/jarviscanada/jarvis_data_eng_SumikshTrehan

Cluster Monitor [GitHub]: Remote Host Monitoring Agent is a professional-grade, agent-based infrastructure monitoring system designed to automate the collection of hardware specifications and real-time resource telemetry from distributed Linux servers. By deploying lightweight Bash scripts on individual nodes, the system continuously pushes critical metricssuch as CPU, memory, and disk utilizationto a centralized PostgreSQL database hosted within Docker. This architecture eliminates the need for manual checks by utilizing Cron-based automation and SQL analytics, empowering System Administrators and DevOps Engineers to proactively monitor server health, identify performance bottlenecks, and drive informed scaling decisions through data-driven insights.

Advanced SQL [GitHub]: This project features a suite of SQL scripts designed to transform raw relational data into actionable business intelligence within a multi-table environment comprising members, facilities, and bookings. The implementation utilizes advanced relational logic, including multi-table joins and correlated subqueries, to analyze facility utilization and member behavior patterns. By leveraging analytical window functions and complex aggregations, the project generates high-level reports on revenue distribution, membership growth, and resource allocation. These queries provide a robust framework for handling real-world data engineering tasks, focusing on query optimization and the extraction of precise, structured datasets for financial and operational reporting.

Python Data Analytics [GitHub]: This project utilizes data analytics and wrangling techniques on a dataset of over one million rows to drive revenue growth through behavioral segmentation. By monitoring key metrics such as monthly sales, cancellation rates, and user acquisition (new vs. existing), the analysis provides a comprehensive overview of business performance. The core of the project is the RFM (Recency, Frequency, Monetary) Framework, which categorizes customers based on their historical transaction patterns to enable precision marketing. By calculating the time since the last purchase, the volume of unique transactions, and total revenue generated per user, the project identifies high-value Whales, loyal Habitual Buyers, and At-Risk and other customer segments, allowing for data-driven strategies to optimize customer lifetime value.

Core Java [GitHub]: The Java Grep App is a command-line search utility developed in Java 8 that replicates the core functionality of the Linux grep tool by recursively scanning directories for specific text patterns. The project demonstrates a transition from traditional imperative programming to a modern functional approach, utilizing Lambdas and the Stream API to build efficient data processing pipelines. By integrating the SLF4J logging framework and Maven for dependency management, the application provides a robust environment for filtering large volumes of text data using Regular Expressions. Designed with scalability in mind, it addresses common performance bottlenecks like memory management through lazy evaluation, ensuring it can handle complex file structures within a Dockerized container environment.

Highlighted Projects

US Population Analysis [GitHub]: This Power BI reporting solution delivers a high-level longitudinal analysis of United States demographic trends from 1970 to 2022. By integrating Azure Maps with OpenStreetMap (OSM) and TomTom layers, the dashboard provides a sophisticated geospatial interface for exploring population distribution across regions, states, and metro areas.

Web Portfolio [GitHub]: Built a modern, theme-aware portfolio using Next.js, Tailwind CSS, and Three.js featuring a chatbot UI for resume-related Q&A and interactive portfolio cards.

Pill Scheduler [GitHub]: Engineered a medication tracker using Next.js 15, Firebase, and Gemini 2.5 Flash for AI visual pill verification while integrating RxCUI identifiers to fetch real-time drug precautions.

Sports Motion Detection and Tracking [GitHub]: Developed a Python and OpenCV pipeline to track motion in sports footage via frame differencing and viewport tracking to generate annotated analysis videos.

Professional Experiences

Data Engineer, Jarvis (Jan 2025-Present): As an Infrastructure and Data Engineer within the Jarvis environment, I was responsible for architecting and maintaining a distributed monitoring ecosystem that bridges the gap between Linux machines and centralized data analytics. Your work focuses on automating hardware inventory and resource tracking across multiple nodes to ensure high availability and operational visibility.

Junior Programmer Co-op, Ontario Public Service (Jan 2025-Apr 2025): As a Junior Programmer Co-op at the Ontario Public Services (TBS Branch) , I engineered a Python automation script using Selenium and Pandas that increased data reporting speed by 85% and saved 20 minutes of manual labor per report. I contributed to system reliability by conducting end-to-end functional testing on over four government tax and dental workflows, resolving logic defects through daily vendor standups. Additionally, I served as a Technical SME, streamlining vendor software alignment by providing critical logic clarifications to ensure all systems remained compliant with Ontario Government program requirements.

Android Developer Co-op, CIBC (May 2024-Aug 2024): At CIBC, as an Android Developer Co-op, I developed and deployed a Kotlin-based gamification feature that boosted user retention by enabling in-app point earning through engagement. I integrated Adobe Analytics to provide real-time user behavior and feature performance insights, driving data-informed product decisions. By managing user stories and defects in JIRA and maintaining comprehensive sprint documentation in Confluence, I accelerated Agile delivery cycles.

Claret Asset Management, IT Specialist (Jun 2022-Jan 2023): As a Junior Programmer Co-op at Ontario Public Services, I engineered Python ETL scripts using Selenium to automate the scraping and storage of over 250 documents into a MySQL database, saving 8+ hours of weekly labor. I modernized legacy infrastructure by migrating SSIS workflows to Python, significantly enhancing system scalability and maintainability. Additionally, I played a key role in facilitating an Azure Cloud migration, collaborating with cross-functional teams and vendors to transition on-premises workloads. I also ensured 100% data integrity for an Angular-based onboarding application by providing backend database support and real-time record synchronization.

Education

Seneca Polytechnic College (2023-2025), Honours Bachelor of Technology - Software Development(BSD), Electrical and Computer Engineering - Scholarship - Dean's List - GPA: 3.8/4.0

Seneca Polytechnic College (2020-2023), Diploma - Computer Programming, Electrical and Computer Engineering - Scholarship - GPA: 3.2/4.0

Georgian College (2015-2017), Diploma - Electrical Engineering, Electrical Engineering - Scholarship - GPA: 3.8/4.0

Miscellaneous

- Linkedin Certifications: Learning Docker, Change Management Foundations, Data Engineering, Learning Power Automate Desktop, Power BI Dashboards, Advanced Snowflake, Apache Kafka, Data Engineering, Azure Spark

Databricks, Learning Apache Airflow, Intro to Spark SQL and DataFrames
- DataCamp Certificates: Object Oriented Programming in Python, Introduction to Python
- Activities/Hobbies: Sports Enthusiast Basketball player, Volunteer